

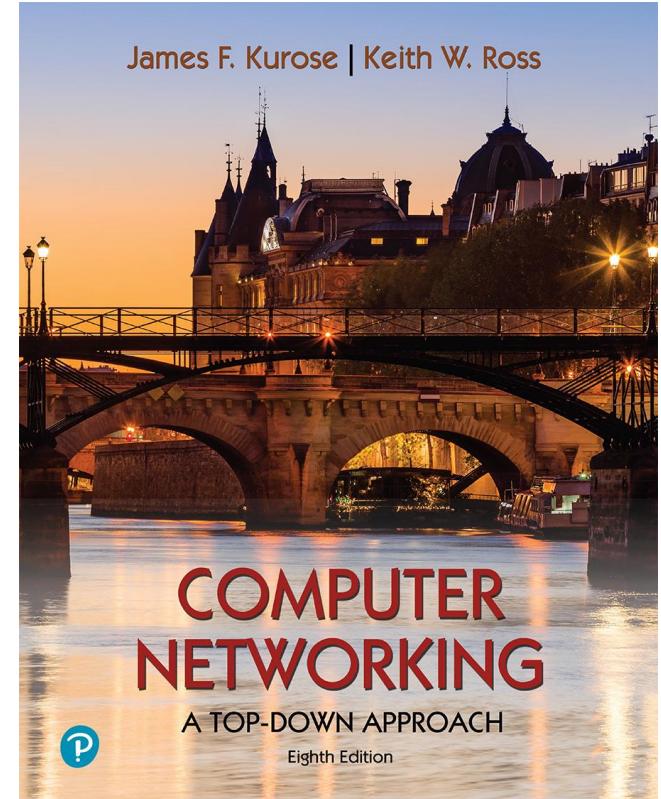
Chapter 2

Application Layer

Yaxiong Xie

Department of Computer Science and Engineering
University at Buffalo, SUNY

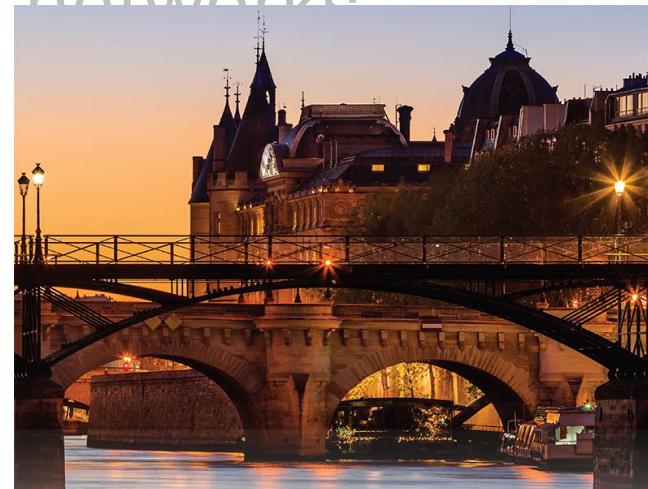
Adapted from the slides of the book's authors



*Computer Networking: A
Top-Down Approach*
8th edition n
Jim Kurose, Keith Ross
Pearson, 2020

Application Layer: Overview

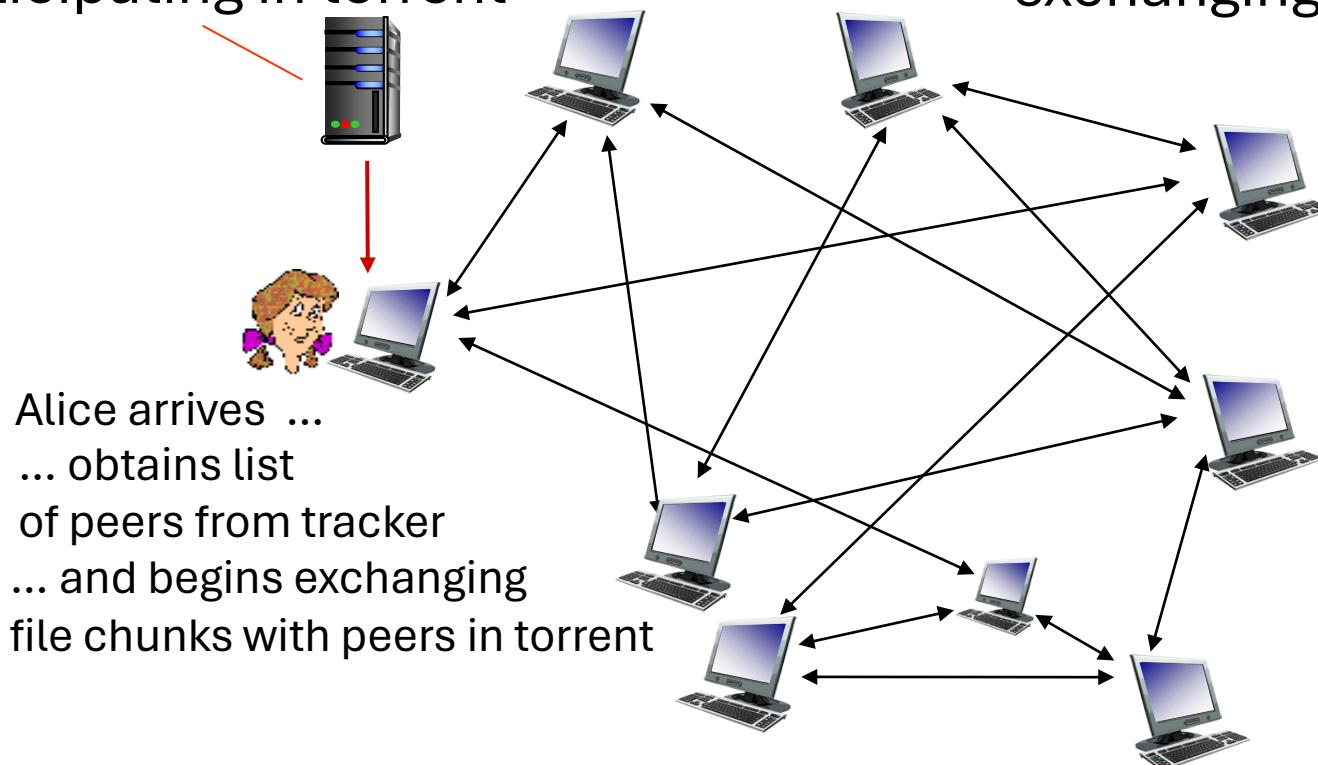
- Principles of network applications
- socket programming with UDP and TCP
- Web and HTTP
- E-mail, SMTP, IMAP
- The Domain Name System DNS
- P2P applications
- video streaming and content distribution networks



P2P file distribution: BitTorrent

- file divided into 256Kb chunks
- peers in torrent send/receive file chunks

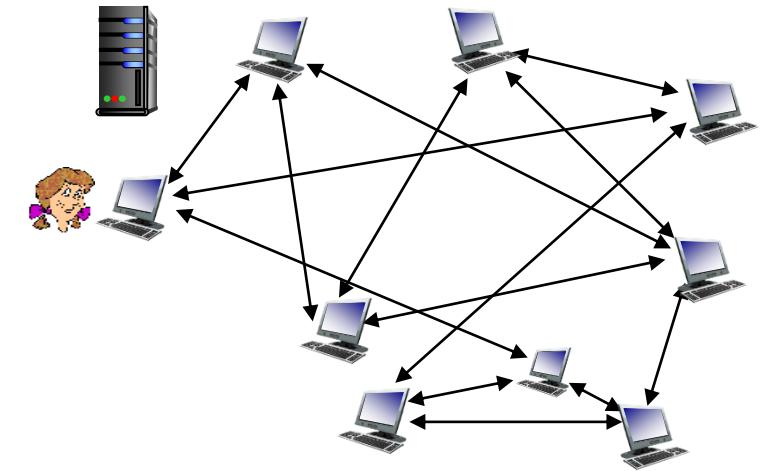
tracker: tracks peers
participating in torrent



torrent: group of peers
exchanging chunks of a file

P2P file distribution: BitTorrent

- peer joining torrent:
 - has no chunks, but will accumulate them over time from other peers
 - registers with tracker to get list of peers, connects to subset of peers (“neighbors”)
- while downloading, peer uploads chunks to other peers
- peer may change peers with whom it exchanges chunks
- *churn*: peers may come and go
- once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent



BitTorrent: requesting, sending file chunks

Requesting chunks:

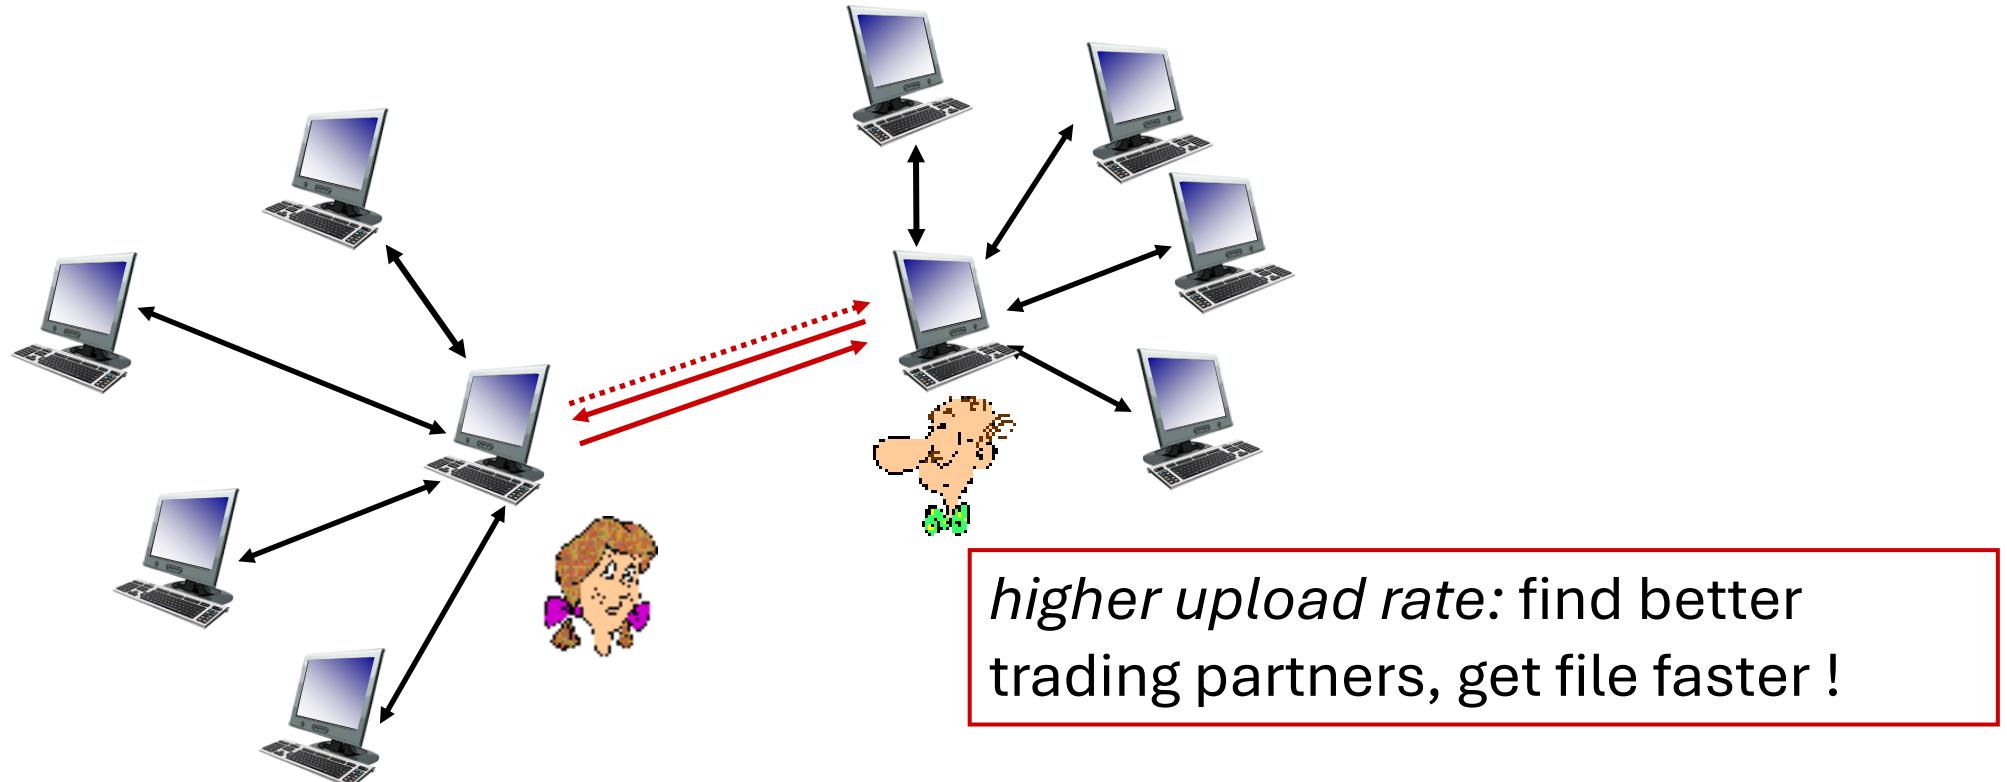
- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each peer for list of chunks that they have
- Alice requests missing chunks from peers, rarest first

Sending chunks: tit-for-tat

- Alice sends chunks to those four peers currently sending her chunks *at highest rate*
 - other peers are choked by Alice (do not receive chunks from her)
 - re-evaluate top 4 every 10 secs
- every 30 secs: randomly select another peer, starts sending chunks
 - “optimistically unchoke” this peer
 - newly chosen peer may join top 4

BitTorrent: tit-for-tat

- (1) Alice “optimistically unchoke” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



Application layer: overview

- Principles of network applications
- socket programming with UDP and TCP
- Web and HTTP
- E-mail, SMTP, IMAP
- The Domain Name System
DNS
- P2P applications
- **video streaming and content distribution networks**



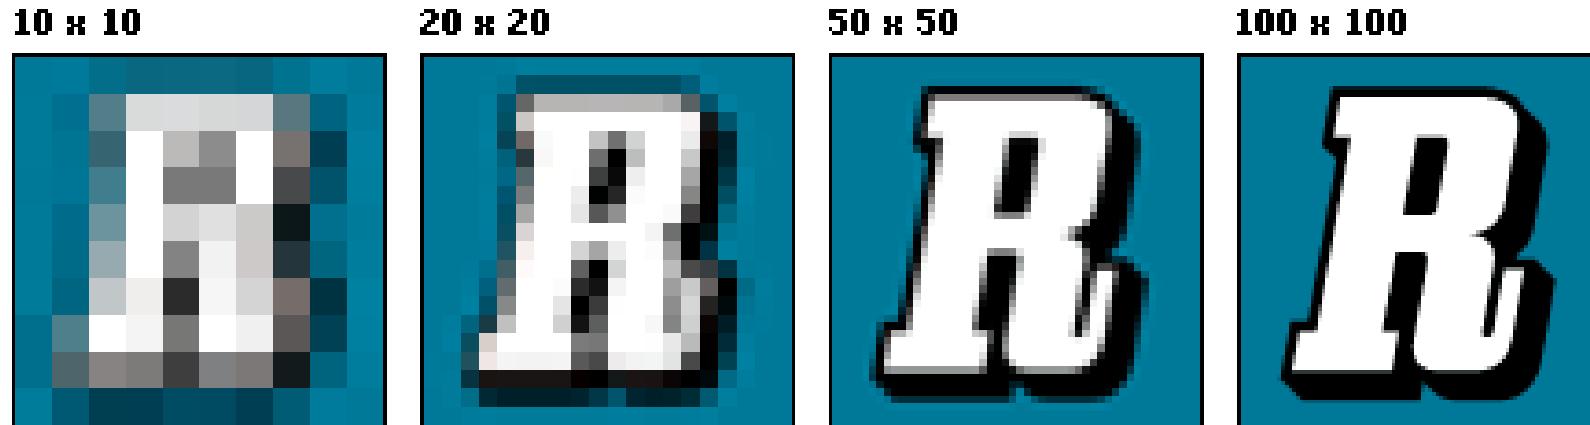
Video Streaming and CDNs: context

- stream video traffic: major consumer of Internet bandwidth
 - Netflix, YouTube, Amazon Prime: 80% of residential ISP traffic (2020)
- *challenge*: scale - how to reach ~1B users?
- *challenge*: heterogeneity
 - different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor)
- *solution*: distributed, application-level infrastructure



Multimedia: video

- video: sequence of images displayed at constant rate
 - e.g., 24 images/sec
- digital image: array of pixels
 - each pixel represented by bits



8000x5000



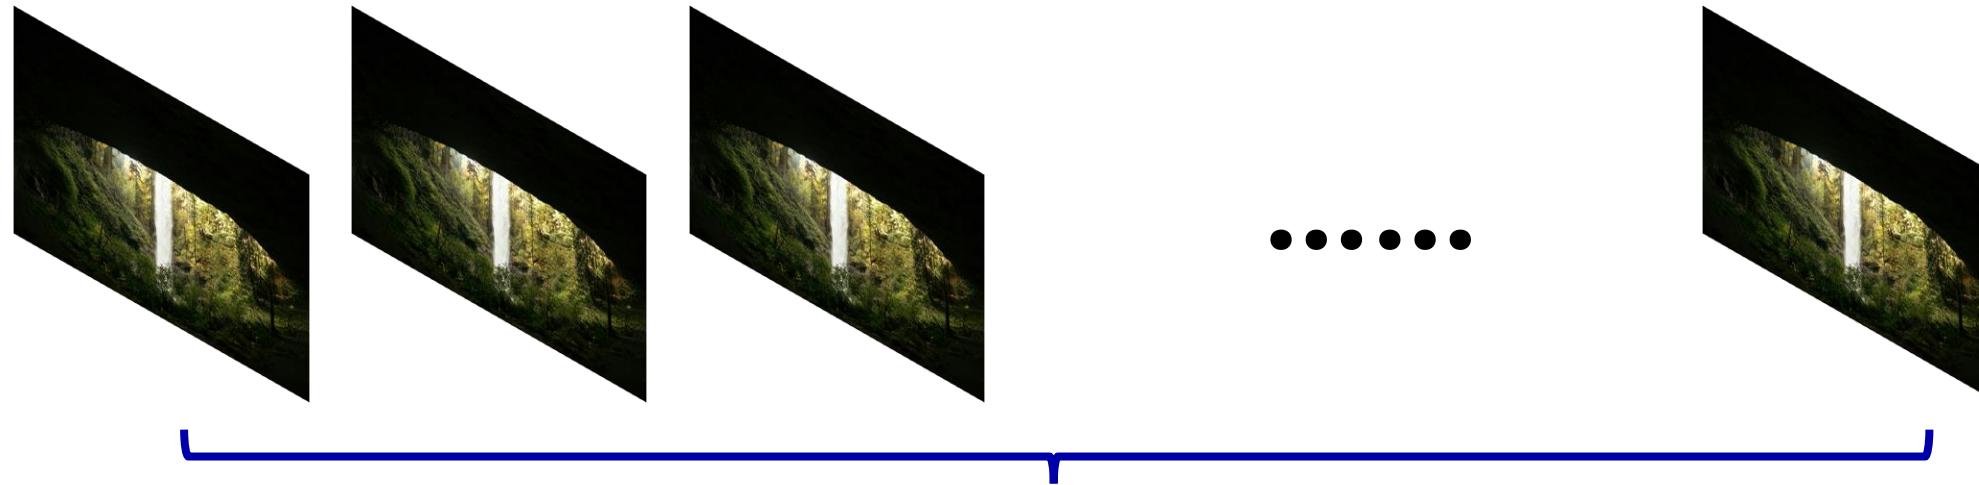
Multimedia: video coding

- Why we need video coding?



Example: a video with **30 fps**

8 MByte = 64 Mbit

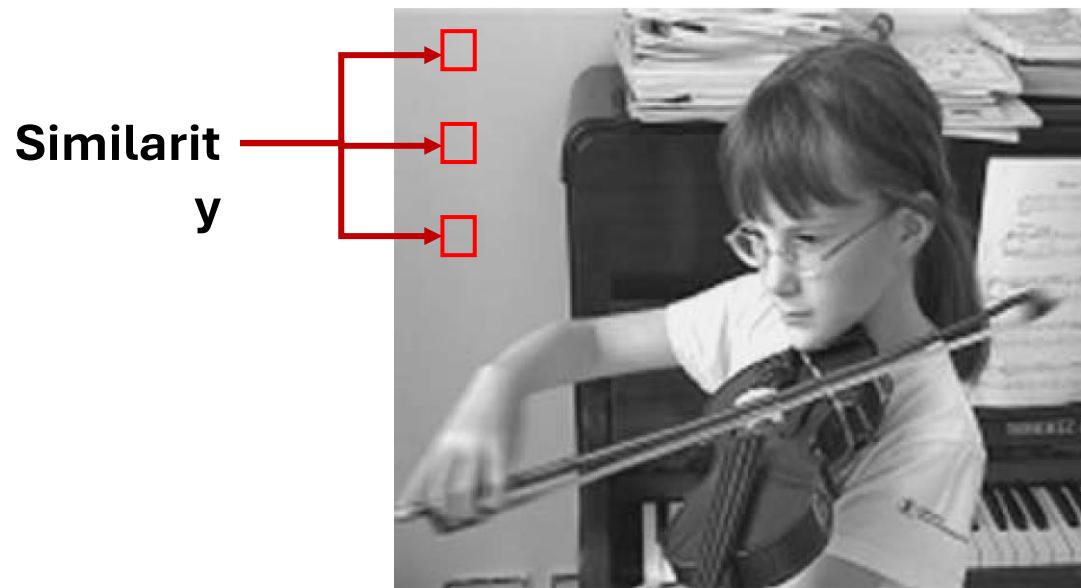


30 images per second

Data Rate: $30 \times 60Mbit = 1800Mbps = 1.8Gbps$

Multimedia: video coding

- coding: leveraging redundancy *within* and *between* images to decrease # bits used to encode image
 - spatial (within image)



Multimedia: video coding

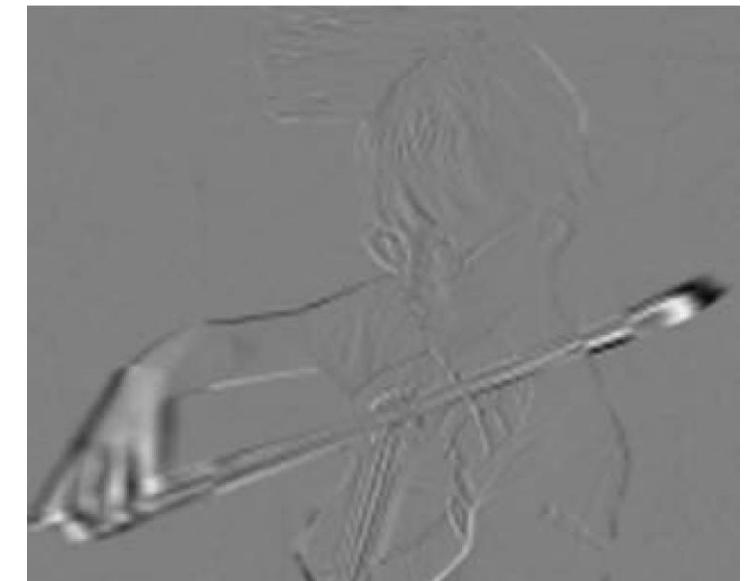
- coding: leveraging **redundancy *within*** and ***between*** images to decrease # bits used to encode image
 - spatial (within image)
 - temporal (from one image to next)



Frame i



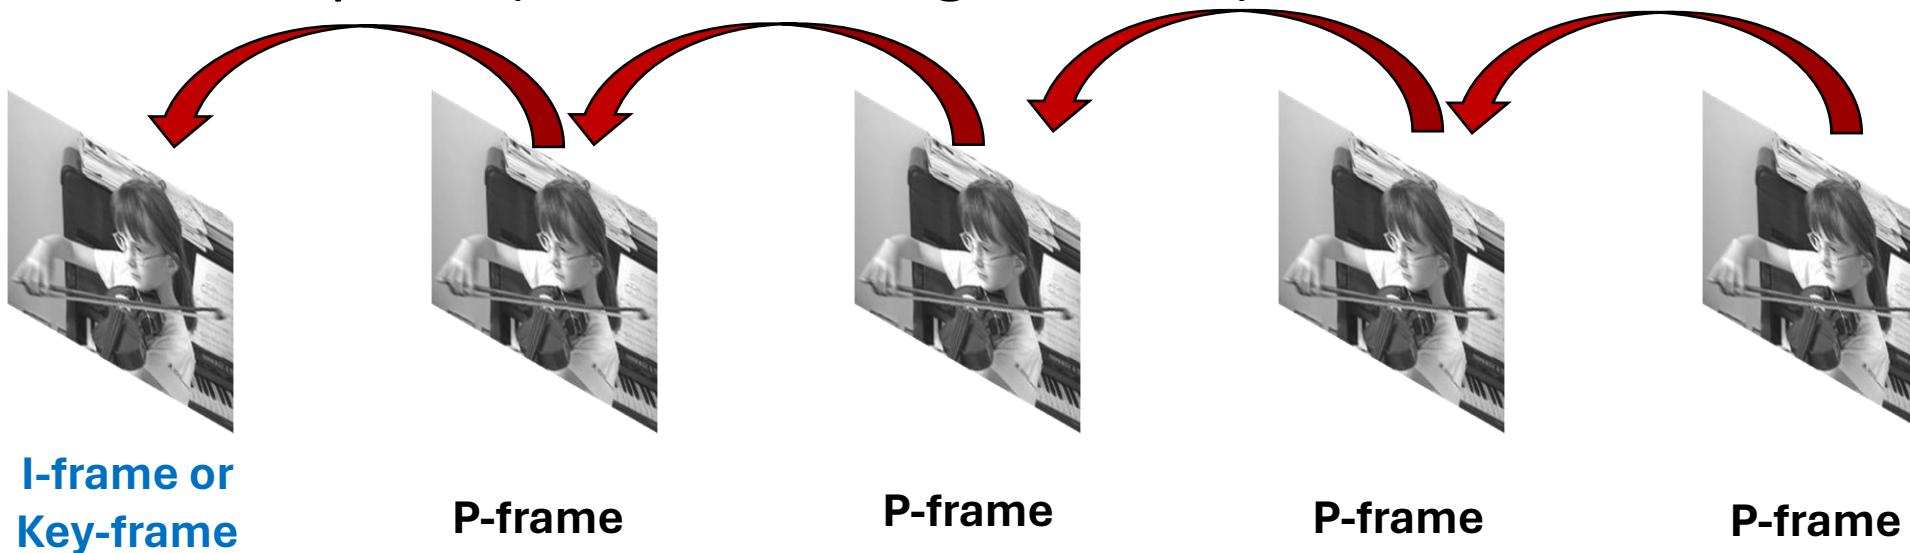
Frame i+1



Difference between frames

Multimedia: video coding

- coding: leveraging **redundancy *within*** and ***between*** images to decrease # bits used to encode image
 - spatial (within image)
 - temporal (from one image to next)

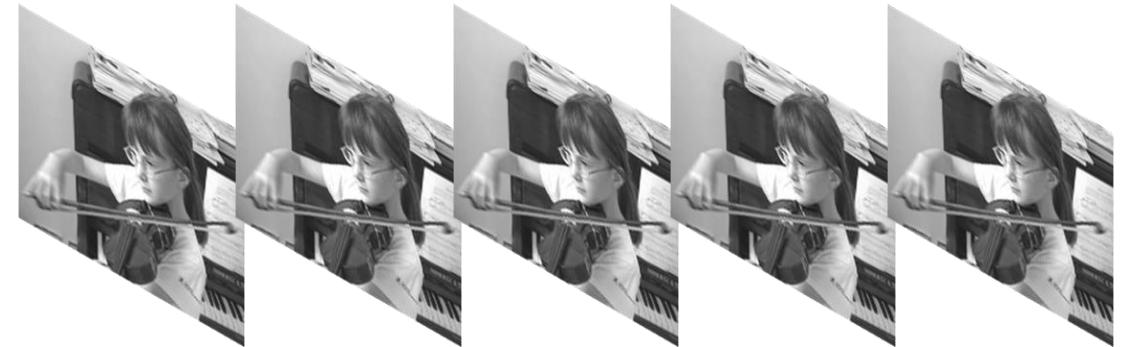


Multimedia: video coding

- coding: leveraging **redundancy *within*** and ***between*** images to decrease # bits used to encode image
 - spatial (within image)
 - temporal (from one image to next)



I-frame or
Key-frame P-frame P-frame P-frame P-frame

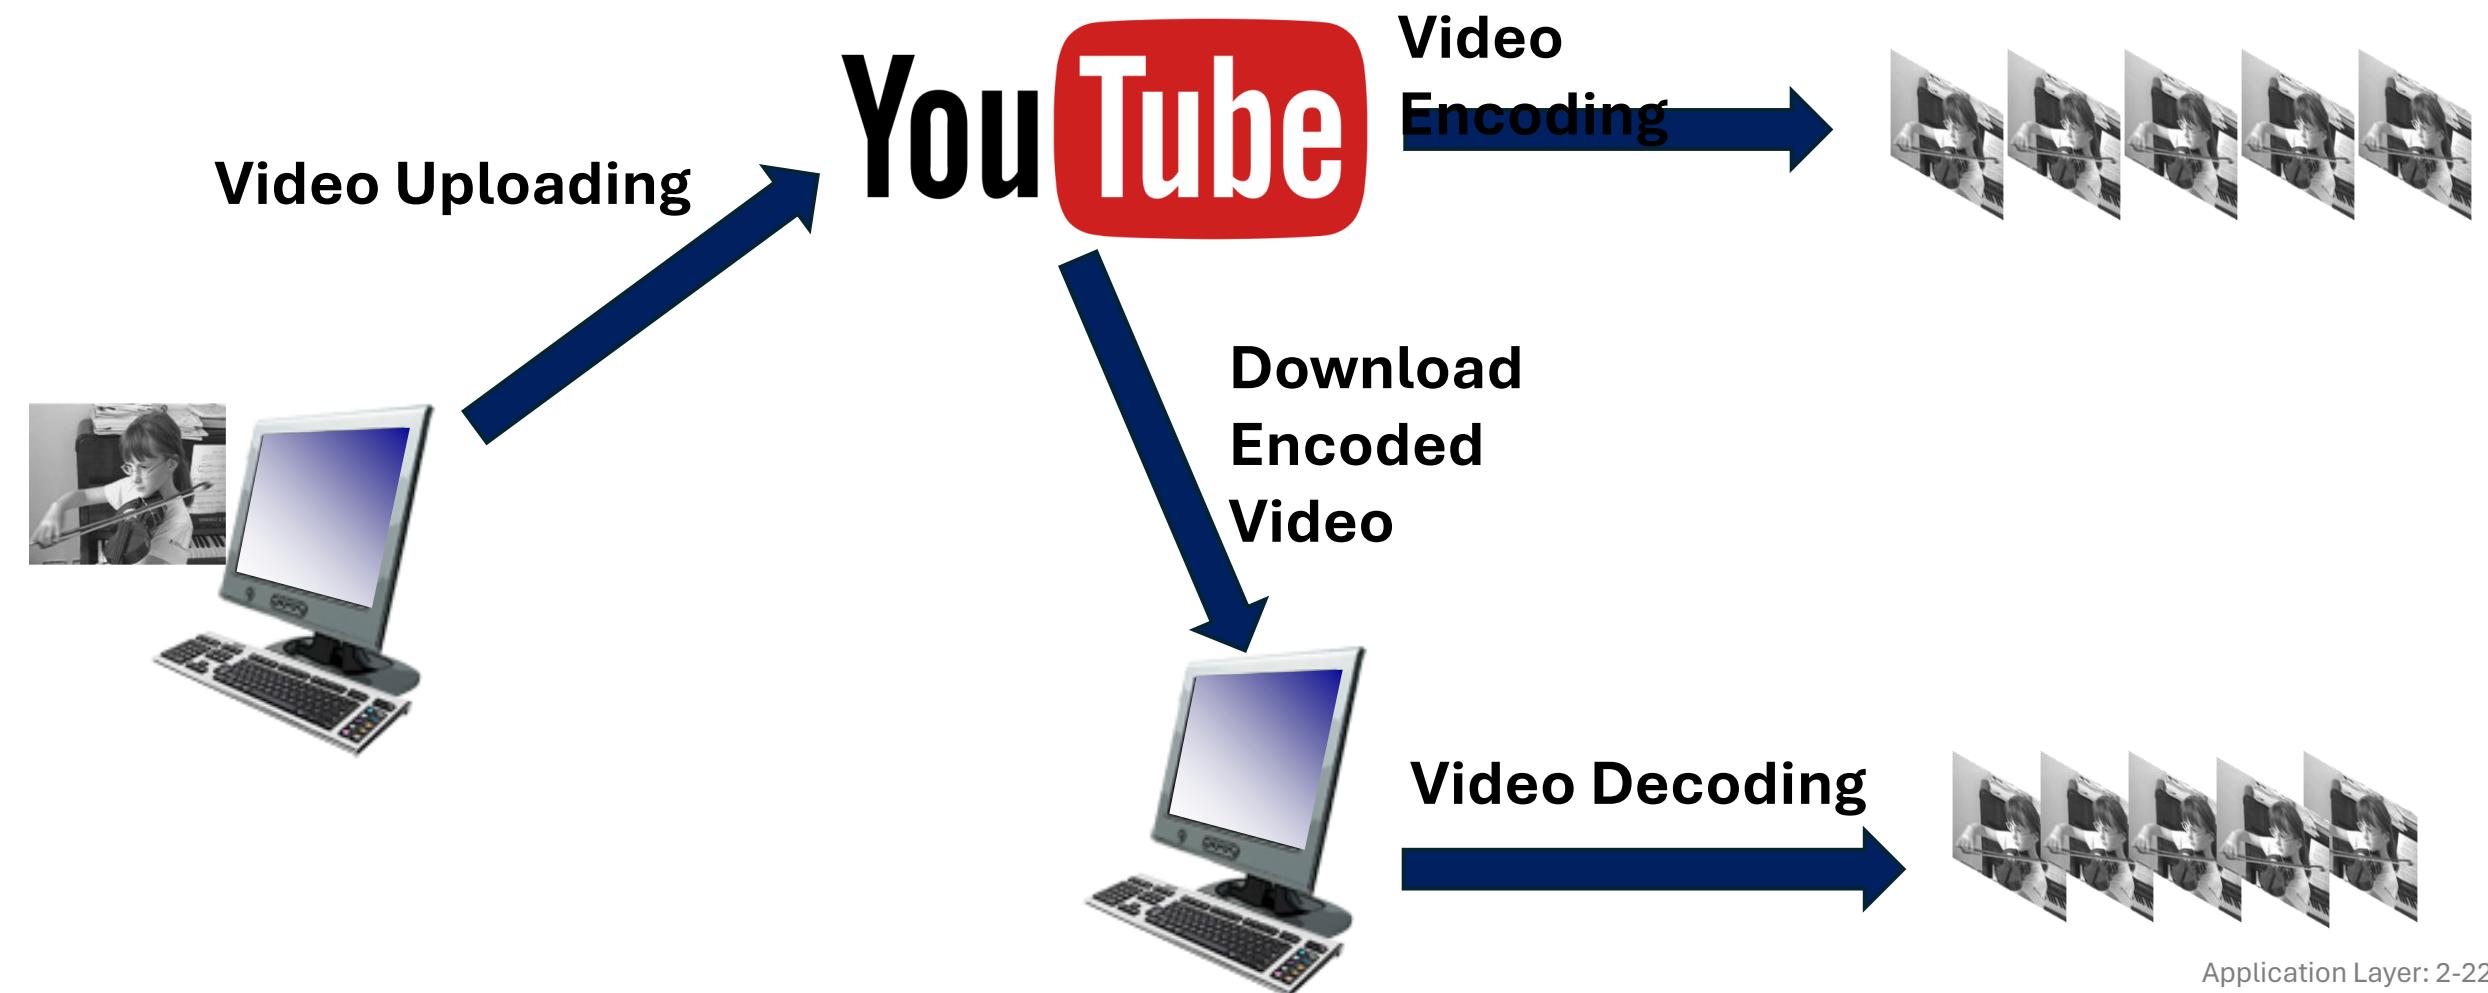


I-frame or
Key-frame P-frame P-frame P-frame P-frame

Video Streaming Applications



Video Streaming Applications: Encoding



Video Streaming Applications: QoE (Quality of User Experience)



Video Start Time

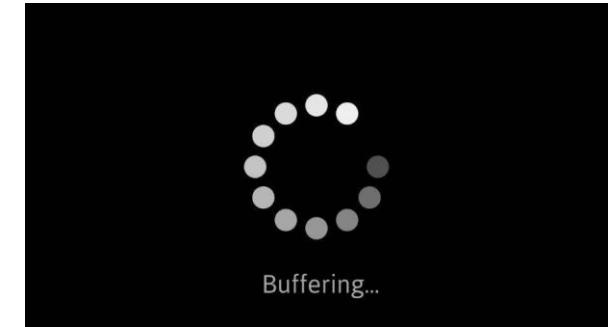
The time between your click of the video and the playing of the video



Video Quality



Rebuffering Event



Streaming multimedia: DASH

*Dynamic, Adaptive
Streaming over HTTP*



Streaming multimedia: DASH

*Dynamic, Adaptive
Streaming over HTTP*



PSY - GENTLEMAN M/V



Subscribe

8.4M Share Download Clip Save ...

1.6B views 10 years ago #PSY #싸이 #GENTLEMAN
PSY - 'I LUV IT' M/V @ PSY - 'I LUV IT' M/V
PSY - 'I LUV IT' M/V - PSY - 'I LUV IT' M/V

Streaming multimedia: DASH

*Dynamic, Adaptive
Streaming over HTTP*



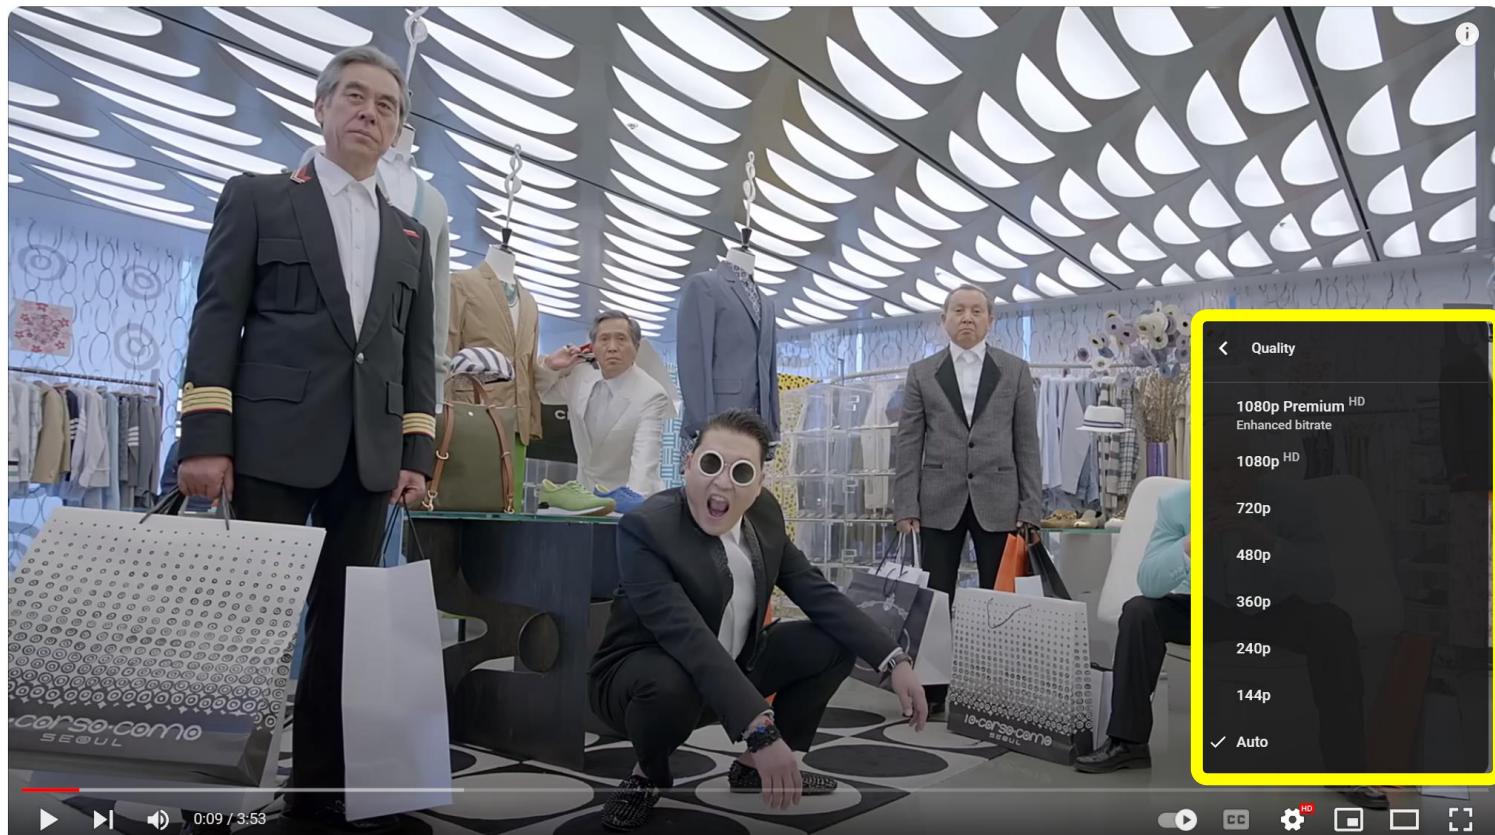
PSY - GENTLEMAN M/V

 officialpsy 18.5M subscribers

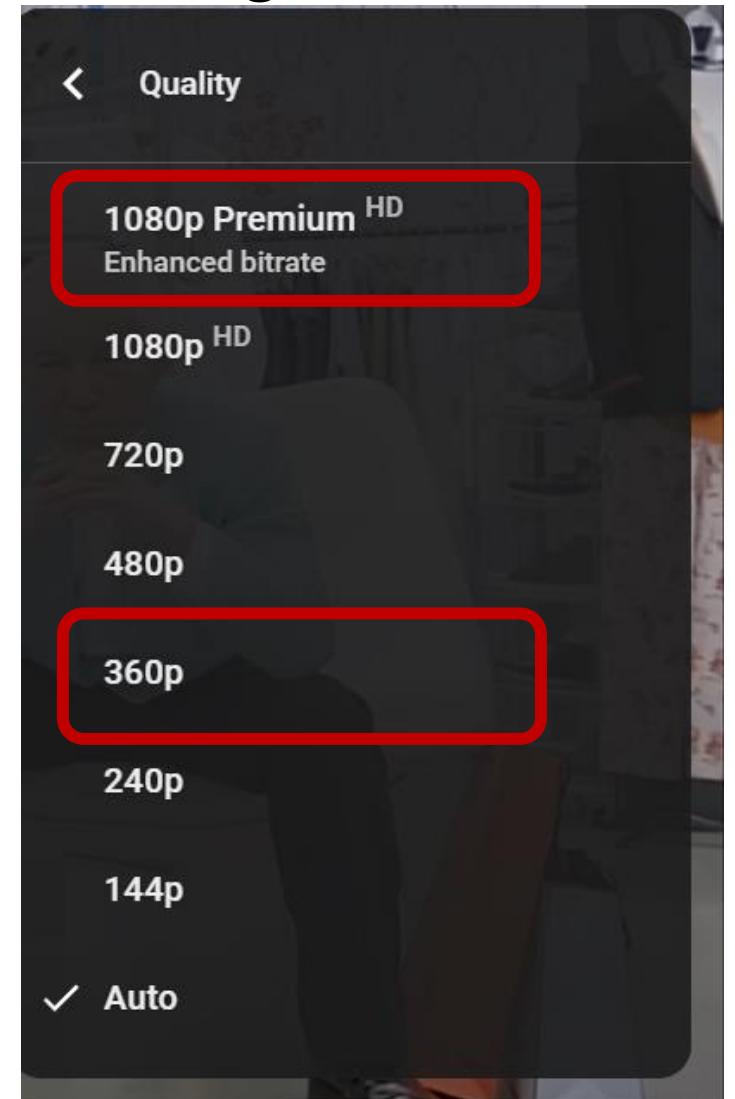
Subscribe

8.4M |  Share  Clip 

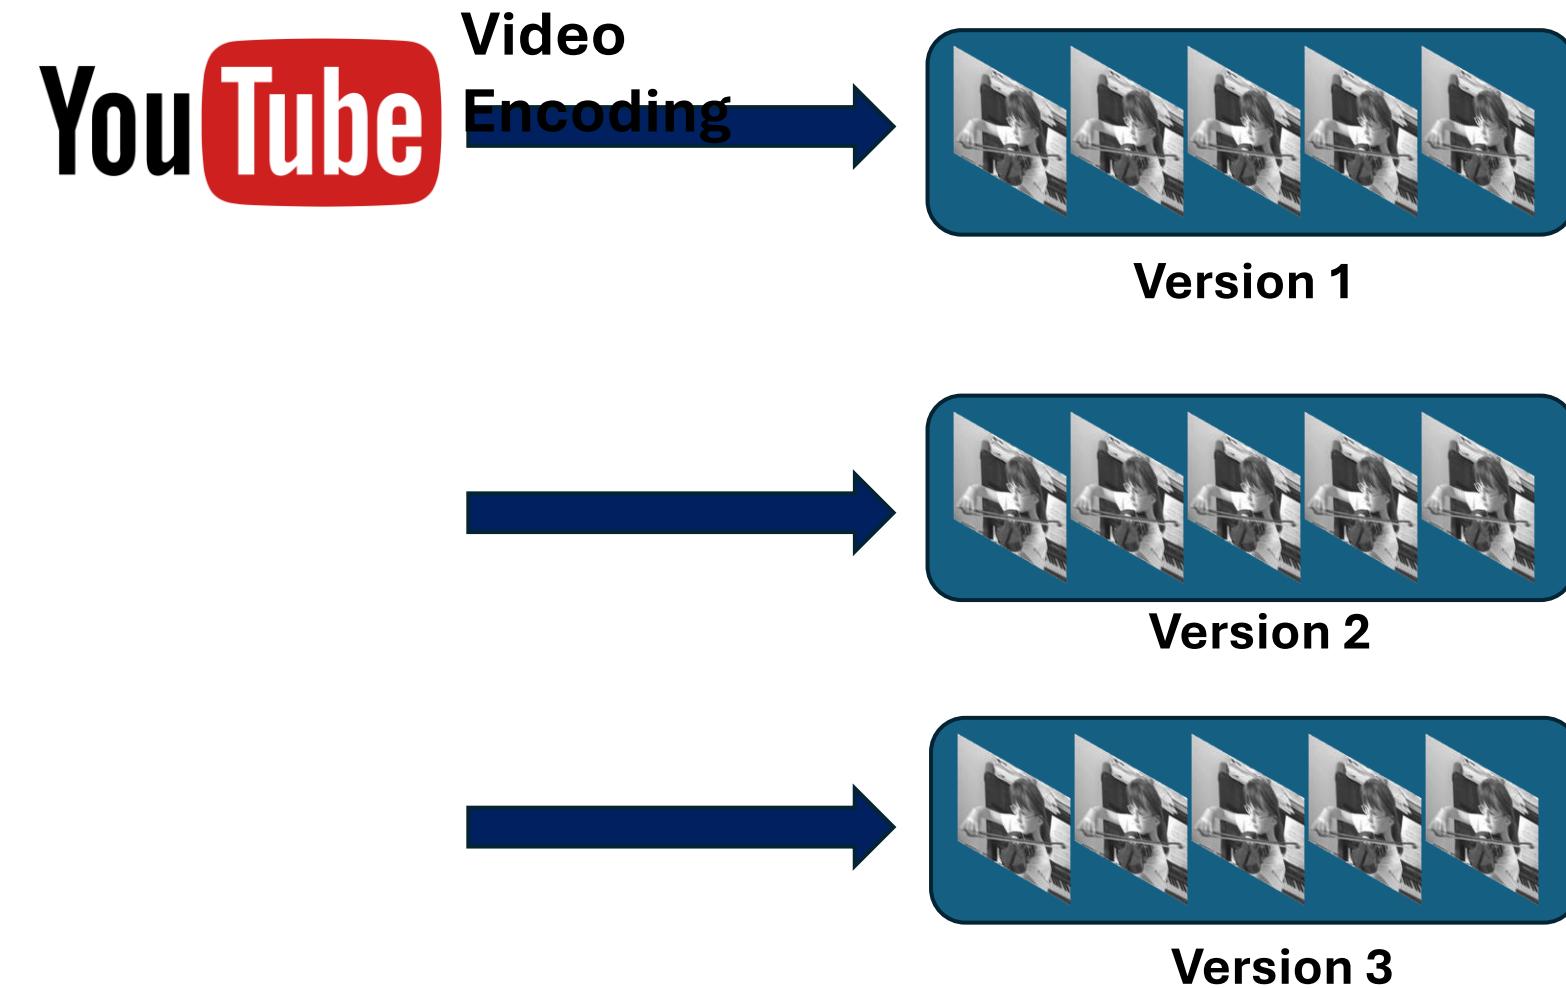
Streaming multimedia: DASH



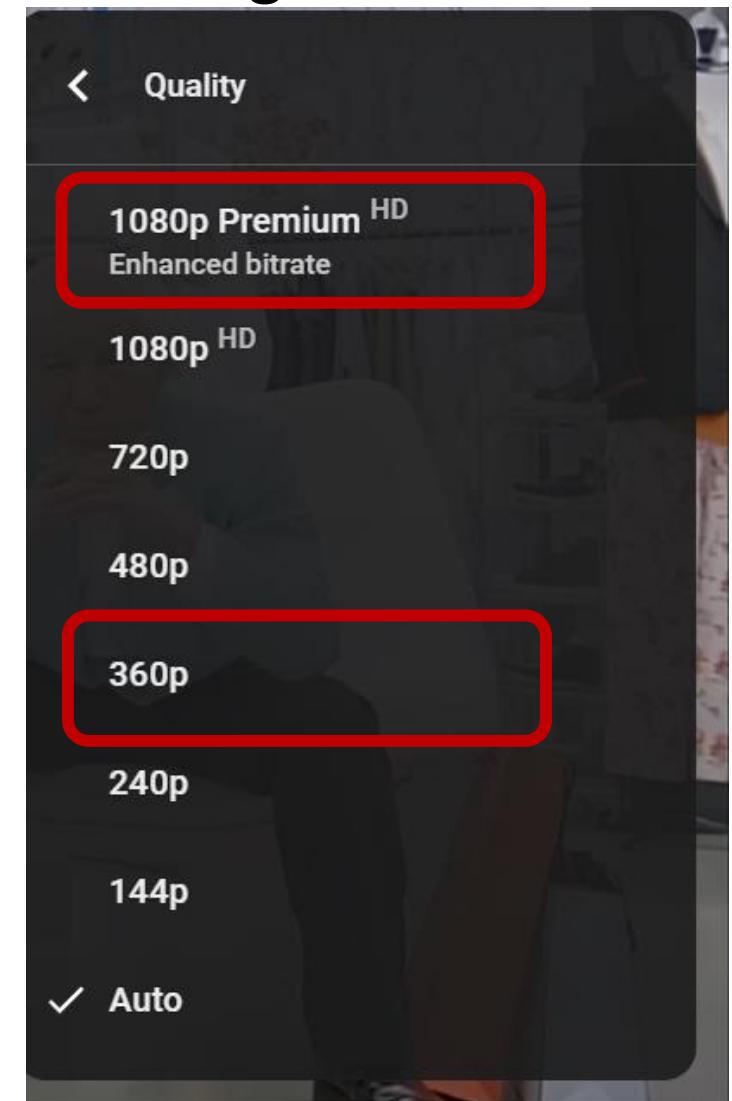
Dynamic, Adaptive Streaming over HTTP



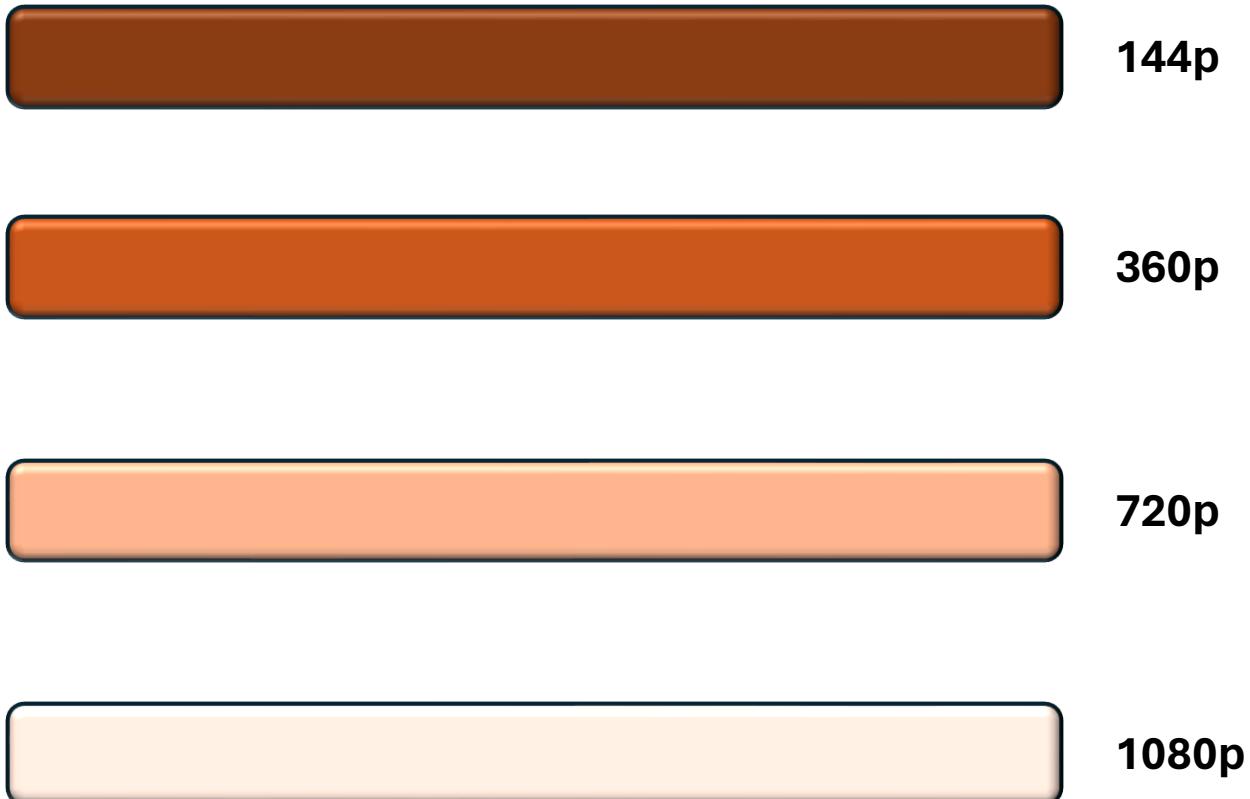
Streaming multimedia: DASH



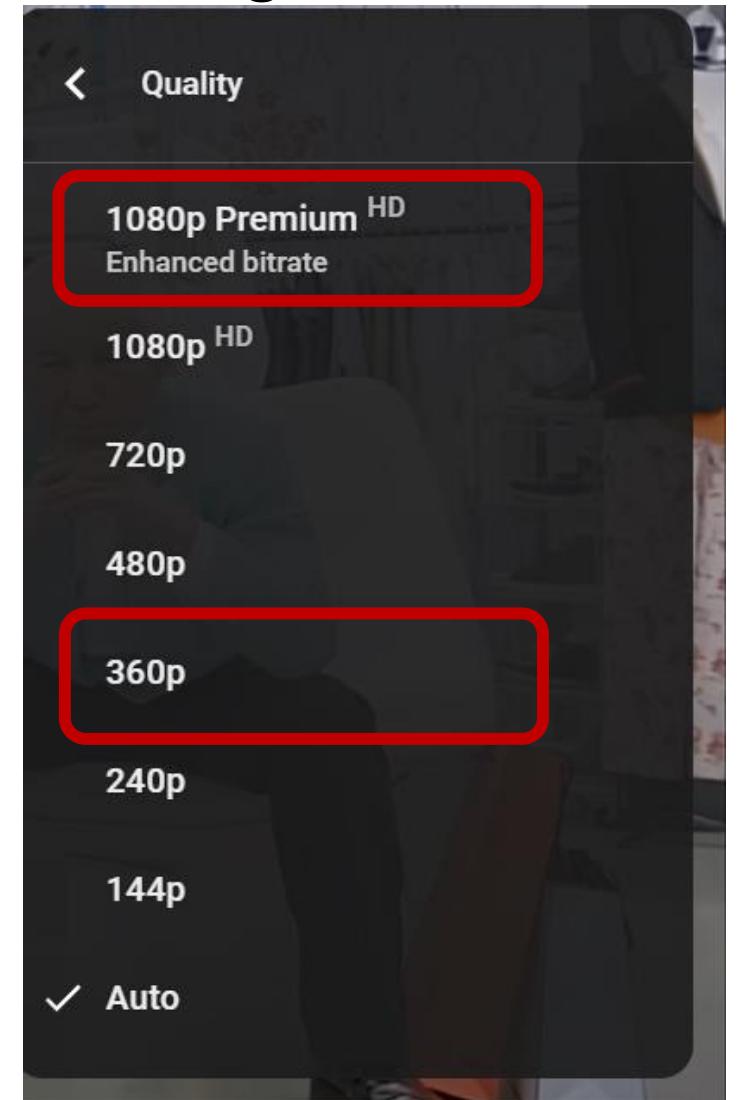
*Dynamic, Adaptive
Streaming over HTTP*



Streaming multimedia: DASH



*Dynamic, Adaptive
Streaming over HTTP*



Streaming multimedia: DASH

Chunk, e.g., 5 seconds



144p



360p

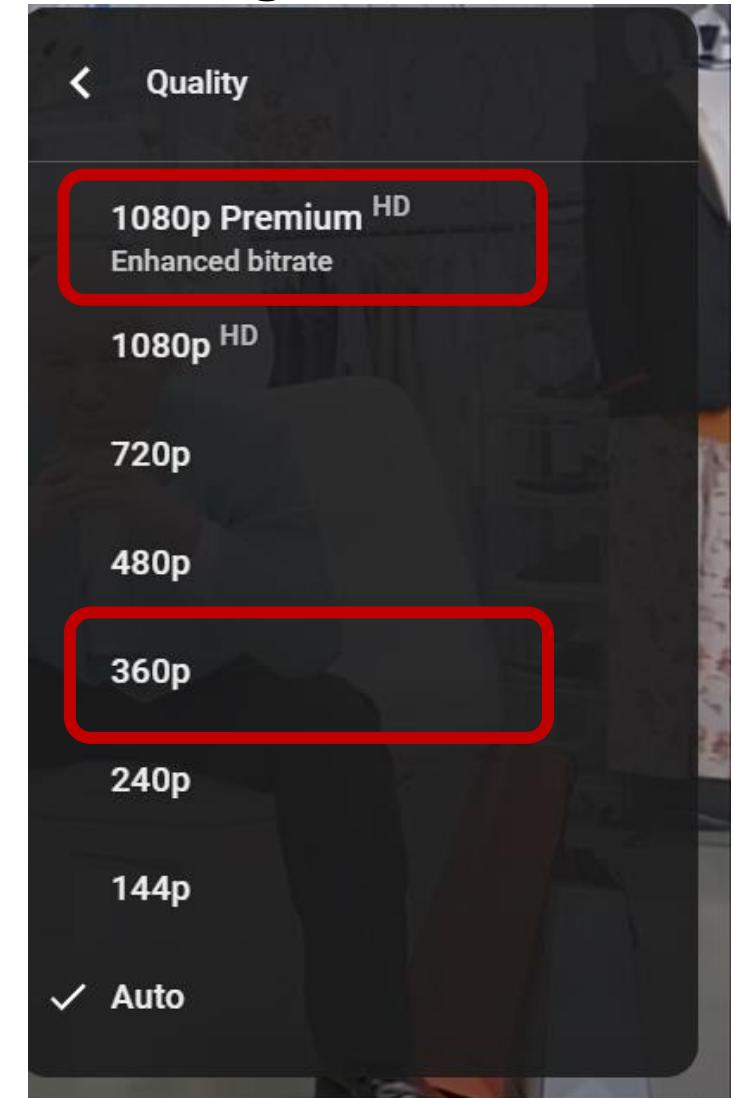


720p



1080p

Dynamic, Adaptive Streaming over HTTP

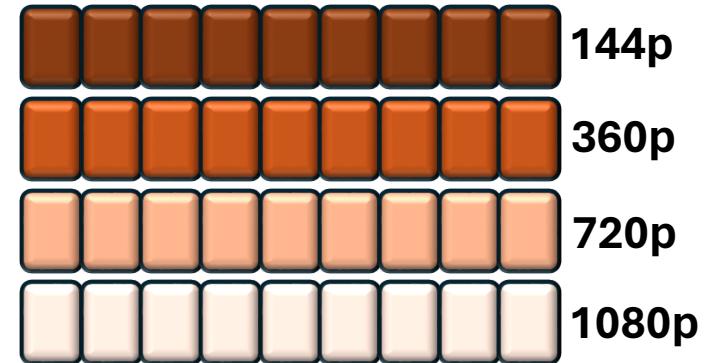


Streaming multimedia: DASH

server:

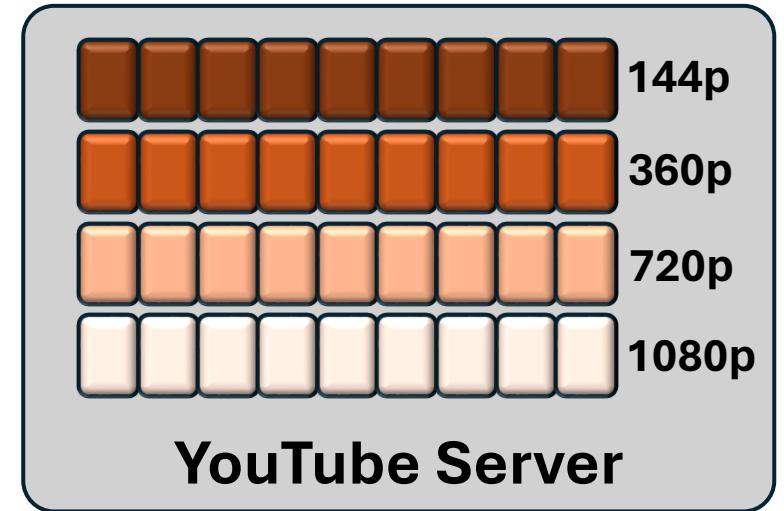
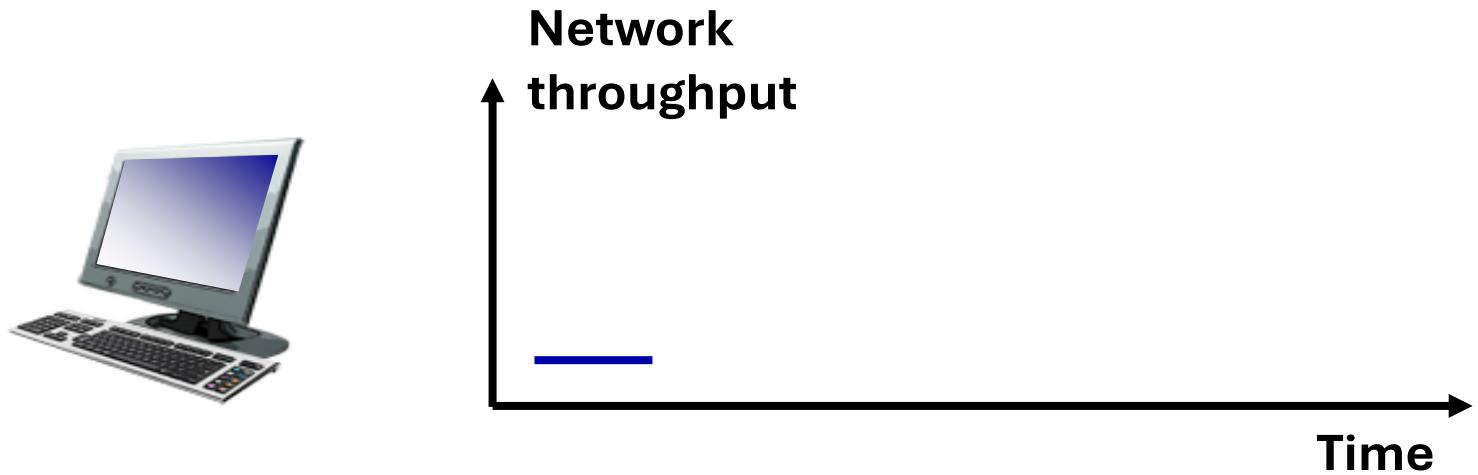
- divides video file into multiple chunks
- each chunk encoded at multiple different rates
- different rate encodings stored in different files
- files replicated in various CDN nodes
- *manifest file*: provides URLs for different chunks

*Dynamic, Adaptive
Streaming over HTTP*



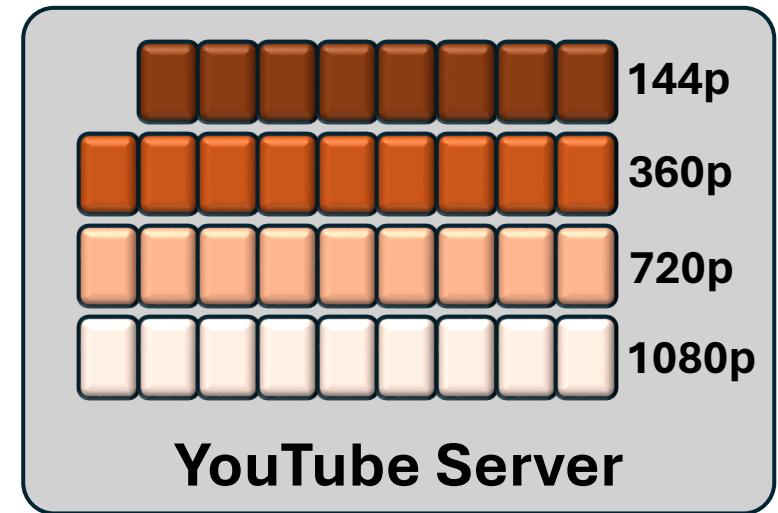
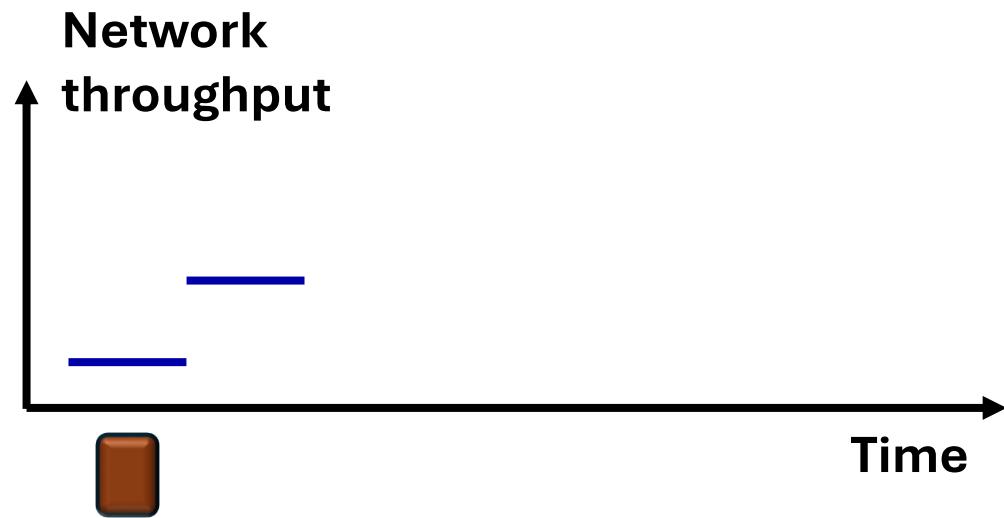
Streaming multimedia: DASH

*Dynamic, Adaptive
Streaming over HTTP*



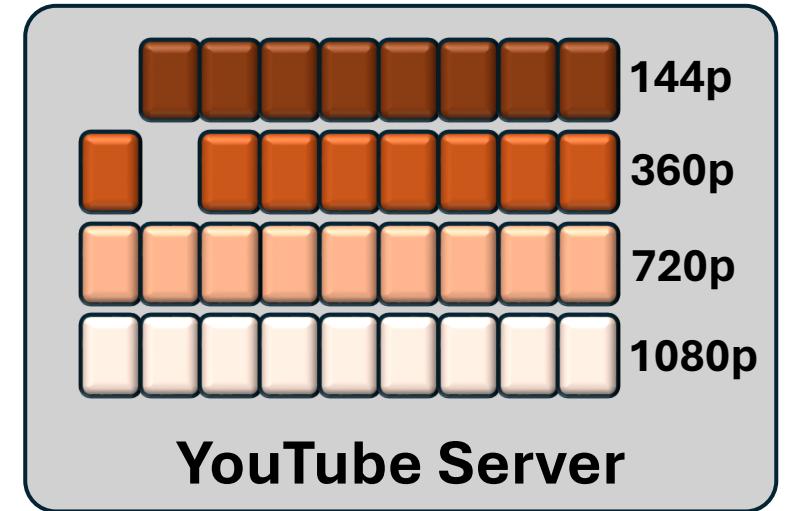
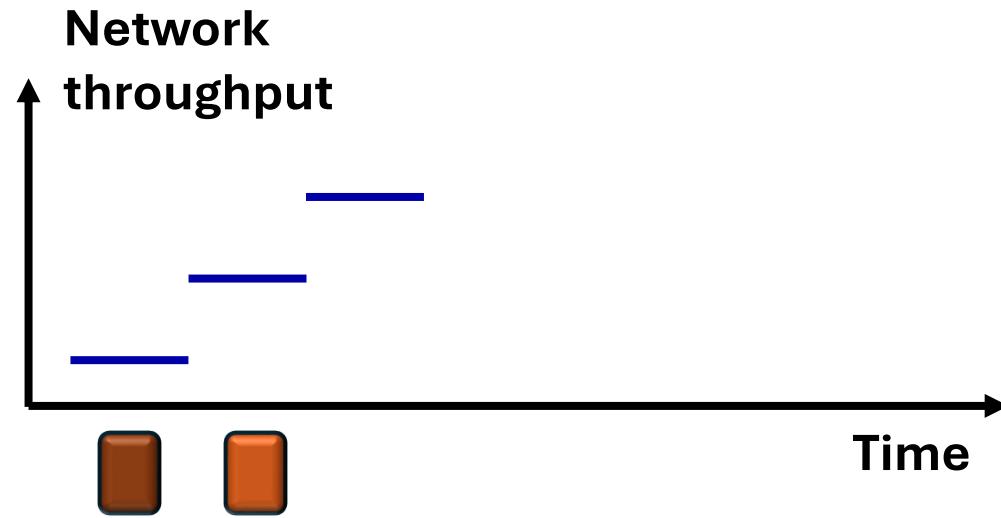
Streaming multimedia: DASH

*Dynamic, Adaptive
Streaming over HTTP*



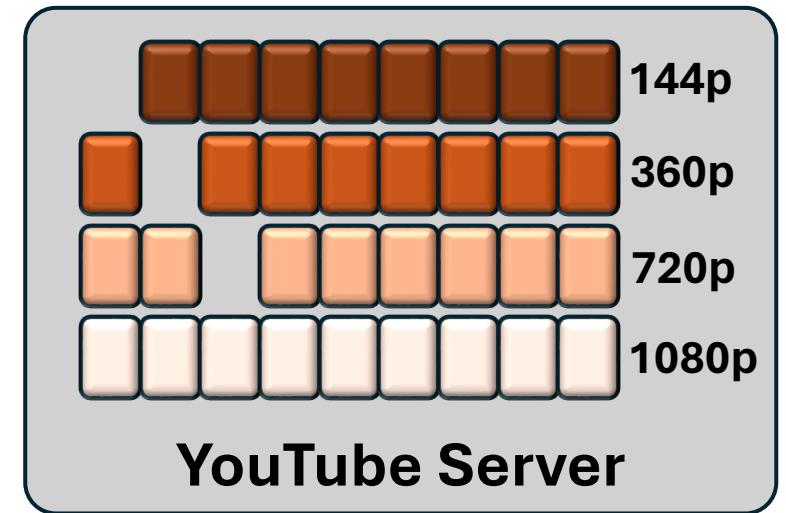
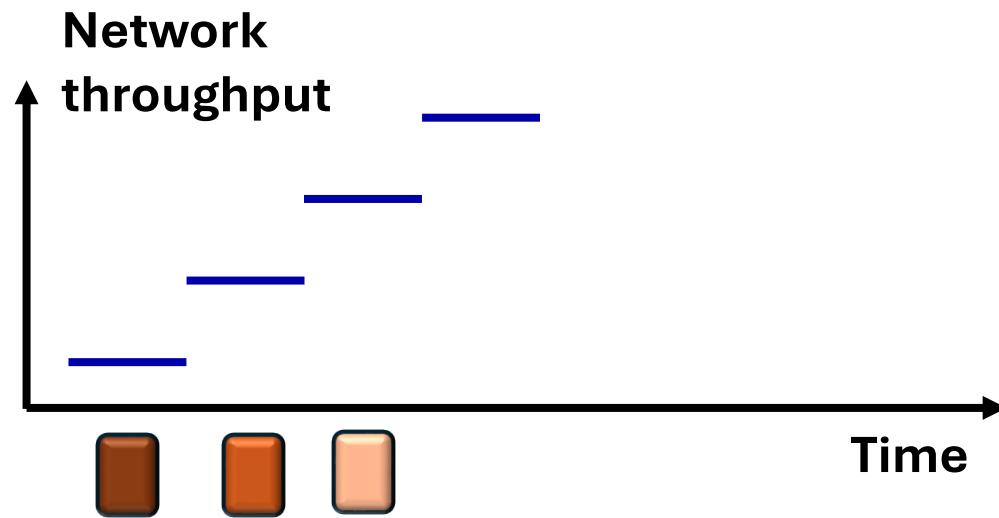
Streaming multimedia: DASH

*Dynamic, Adaptive
Streaming over HTTP*



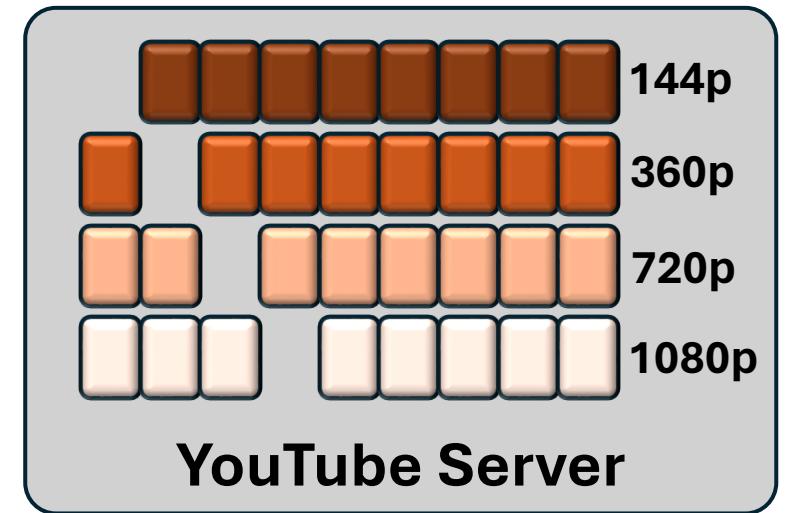
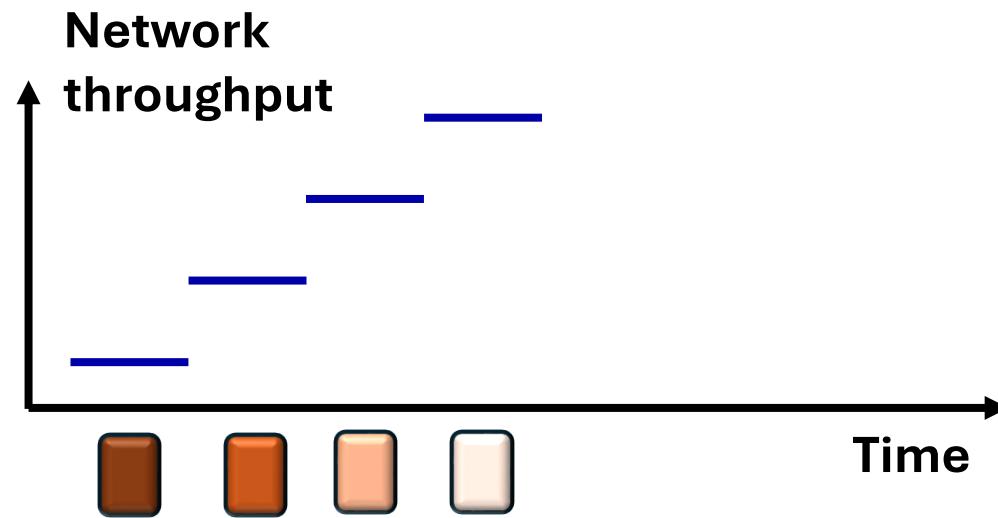
Streaming multimedia: DASH

*Dynamic, Adaptive
Streaming over HTTP*



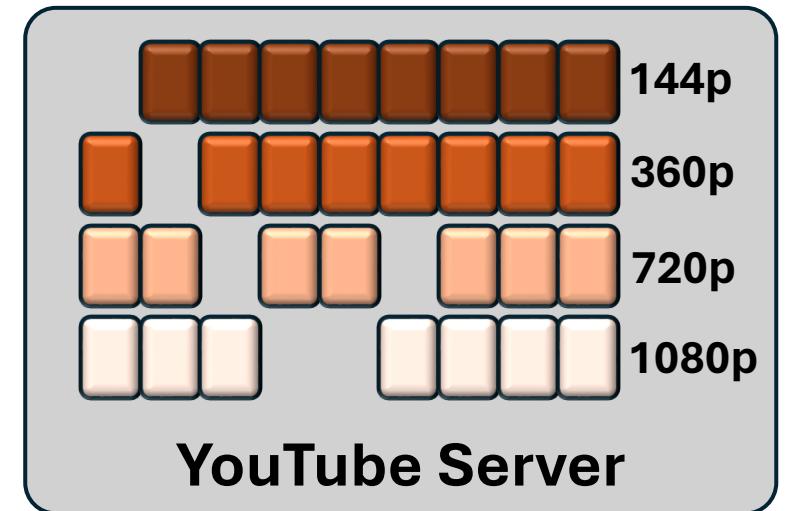
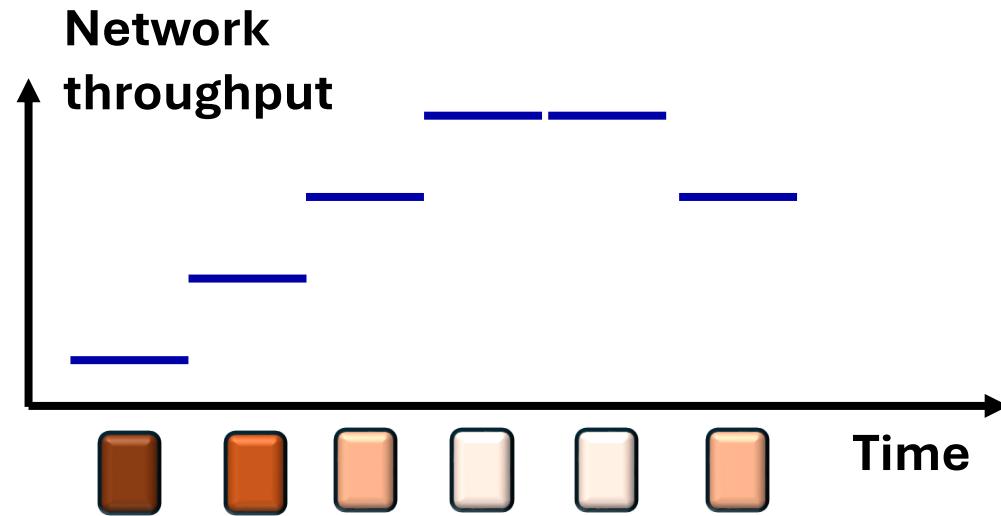
Streaming multimedia: DASH

*Dynamic, Adaptive
Streaming over HTTP*



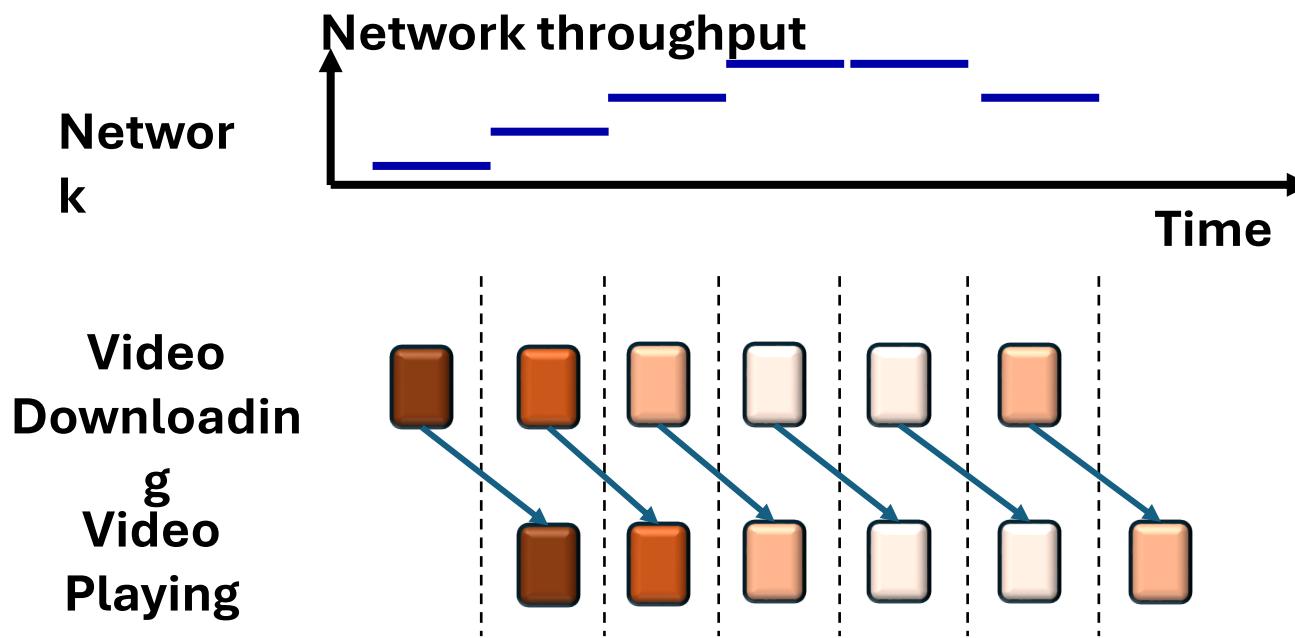
Streaming multimedia: DASH

*Dynamic, Adaptive
Streaming over HTTP*



Streaming multimedia: DASH

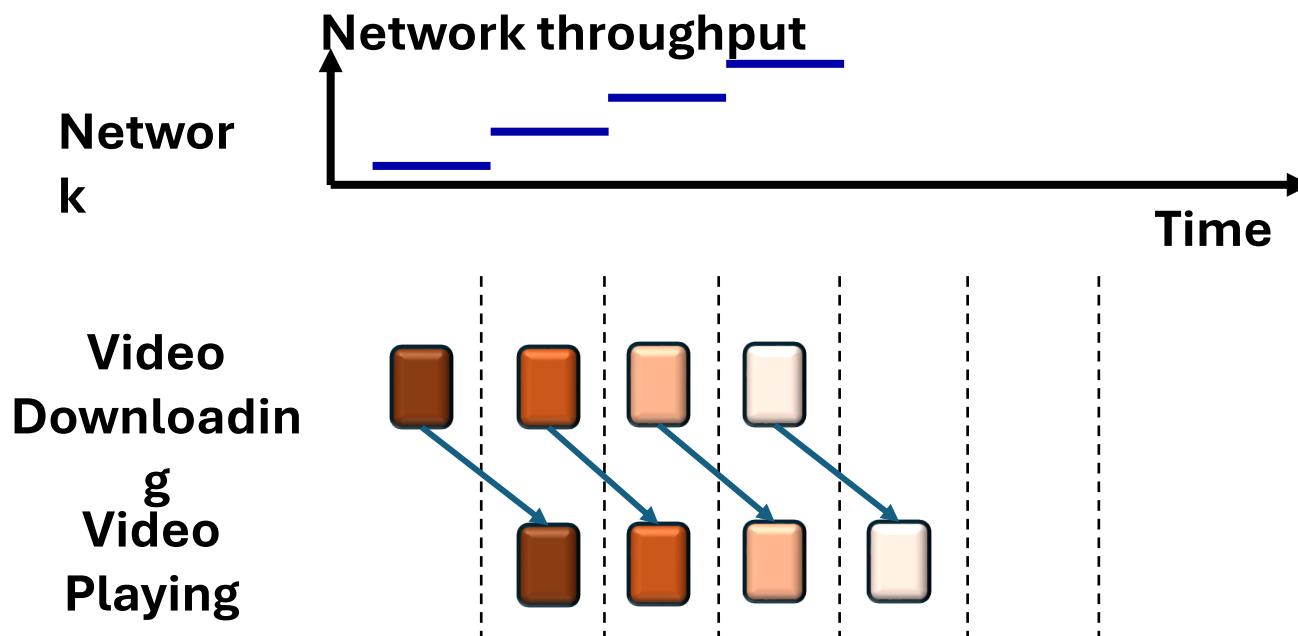
*Dynamic, Adaptive
Streaming over HTTP*



Minimize Video Start Time

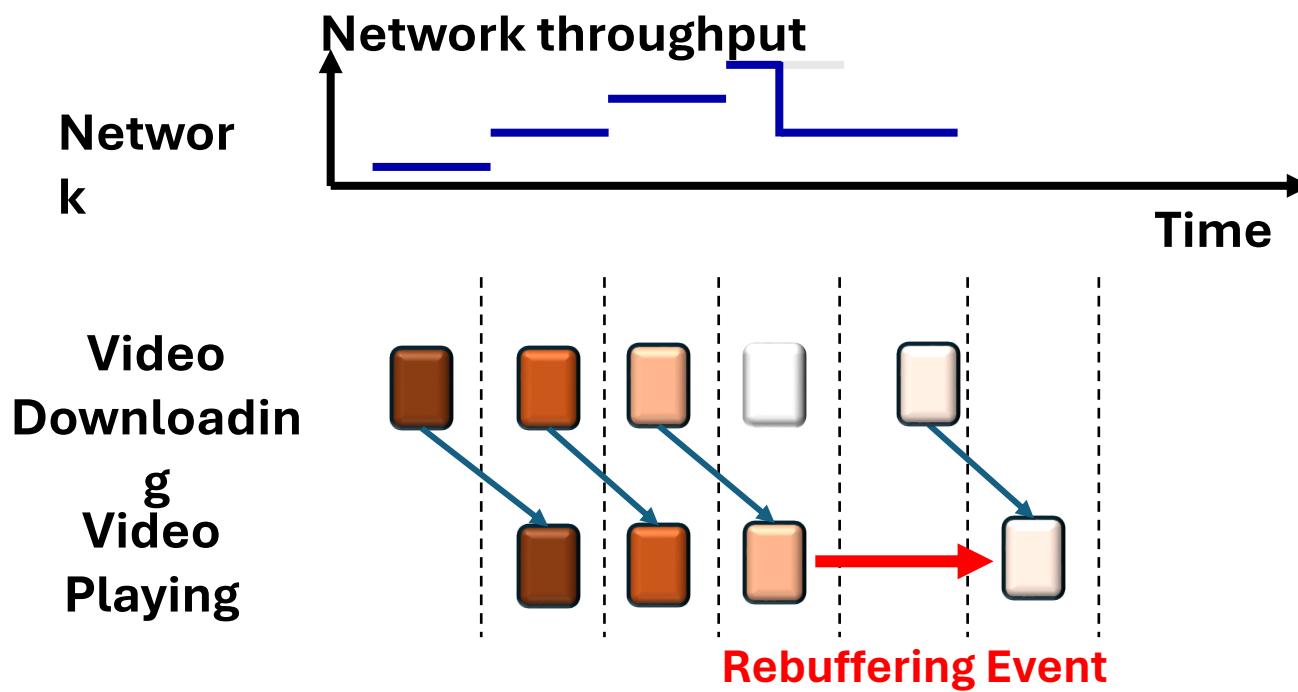
Streaming multimedia: DASH

*Dynamic, Adaptive
Streaming over HTTP*



Streaming multimedia: DASH

*Dynamic, Adaptive
Streaming over HTTP*

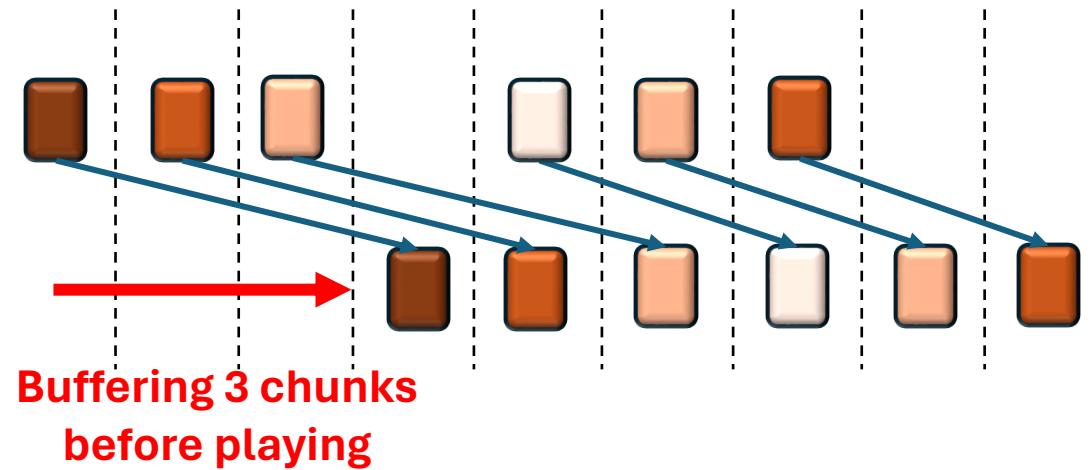


Streaming multimedia: DASH

*Dynamic, Adaptive
Streaming over HTTP*



Video
Downloadin
g
Video
Playing



Streaming multimedia: DASH

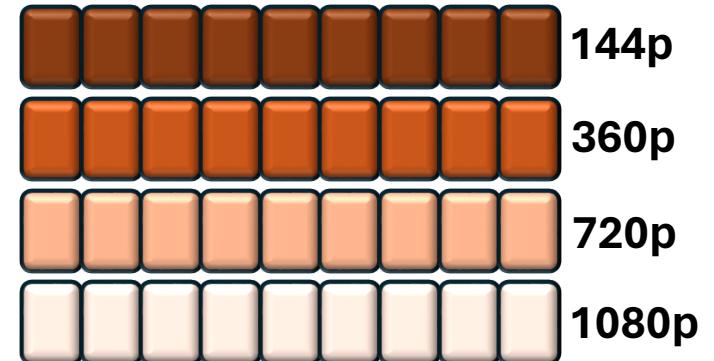
*Dynamic, Adaptive
Streaming over HTTP*

server:

- divides video file into multiple chunks
- each chunk encoded at multiple different rates
- different rate encodings stored in different files
- files replicated in various CDN nodes
- *manifest file*: provides URLs for different chunks

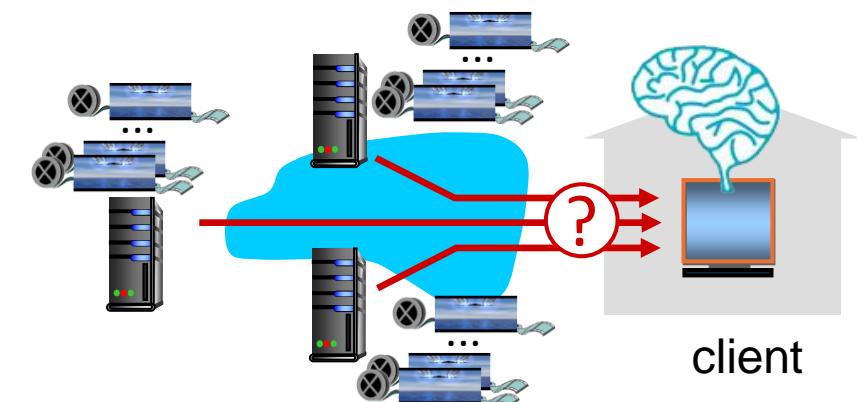
client:

- periodically estimates server-to-client bandwidth
- consulting manifest, requests one chunk at a time
 - chooses maximum coding rate sustainable given current bandwidth
 - can choose different coding rates at different points in time (depending on available bandwidth at time), and from different servers



Streaming multimedia: DASH

- “*intelligence*” at client: client determines
 - *when* to request chunk (so that buffer starvation, or overflow does not occur)
 - *what encoding rate* to request (higher quality when more bandwidth available)
 - *where* to request chunk (can request from URL server that is “close” to client or has high available bandwidth)

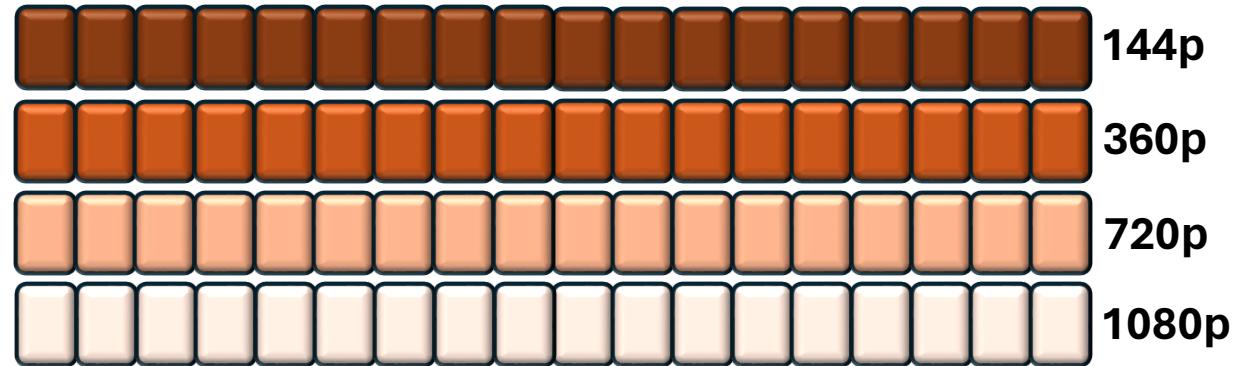


Streaming video = encoding + DASH + playout buffering

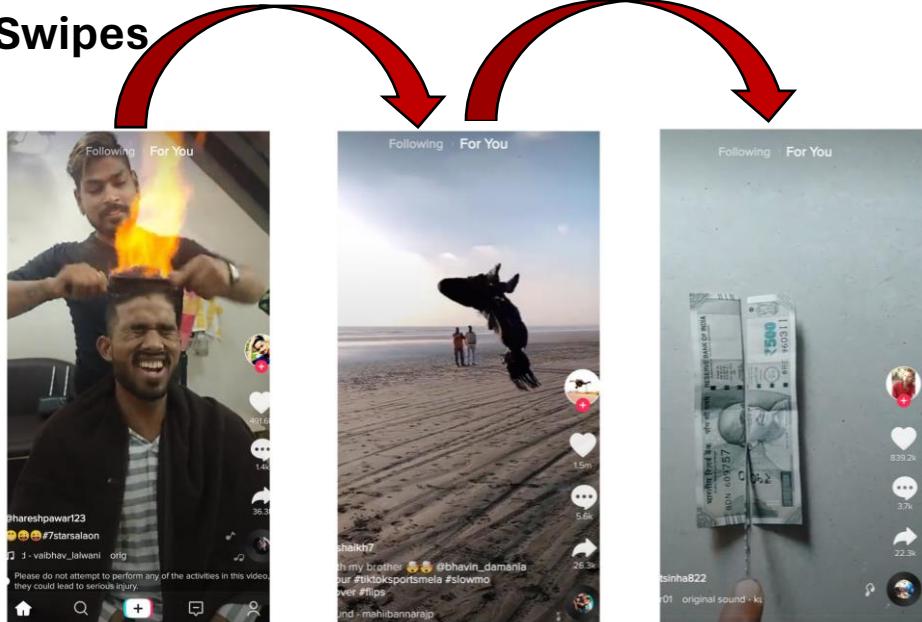
Video Streaming Applications



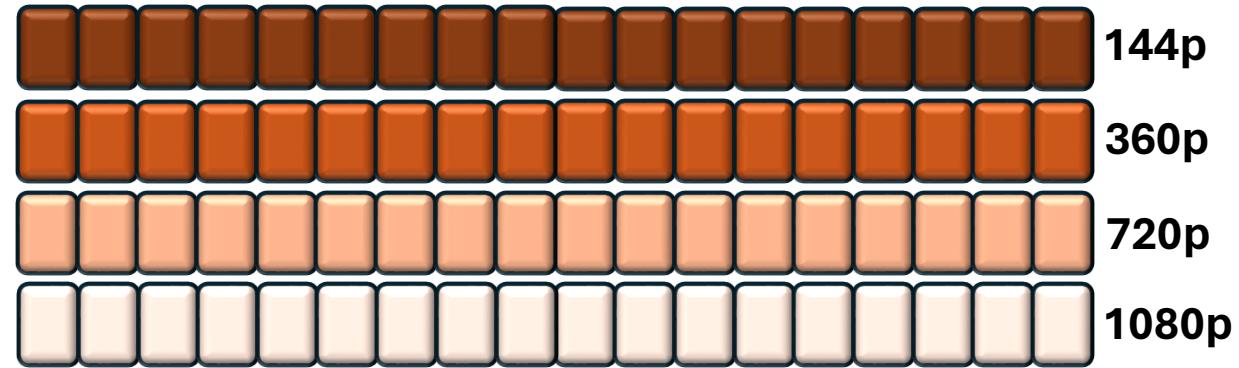
Video Streaming Applications



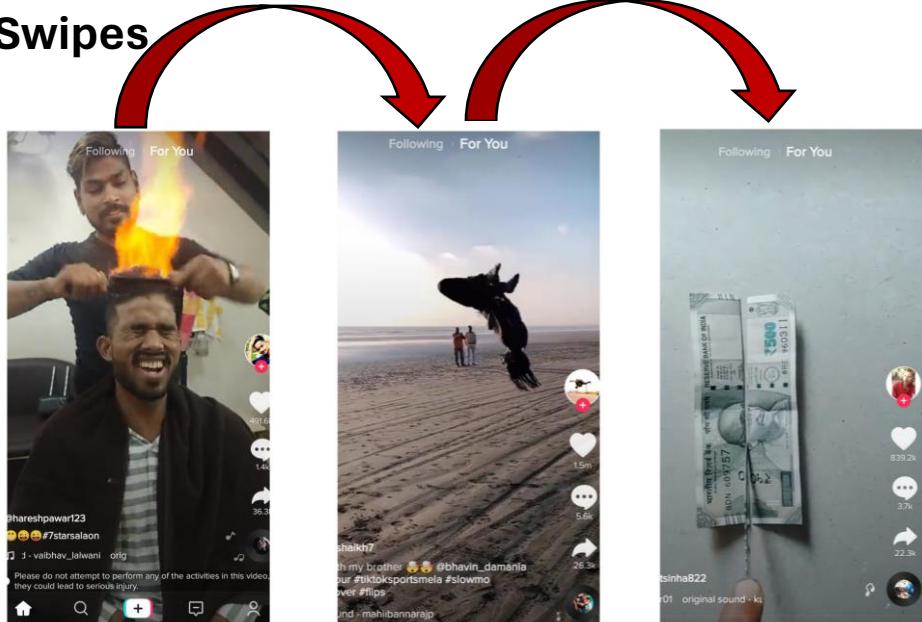
User Swipes



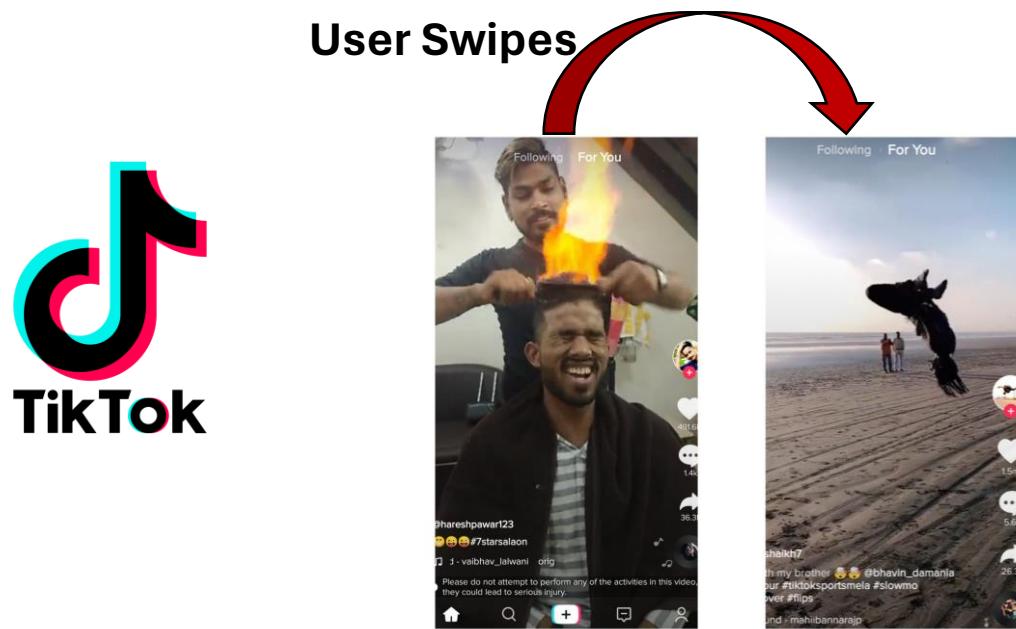
Video Streaming Applications



User Swipes



Video Streaming Applications



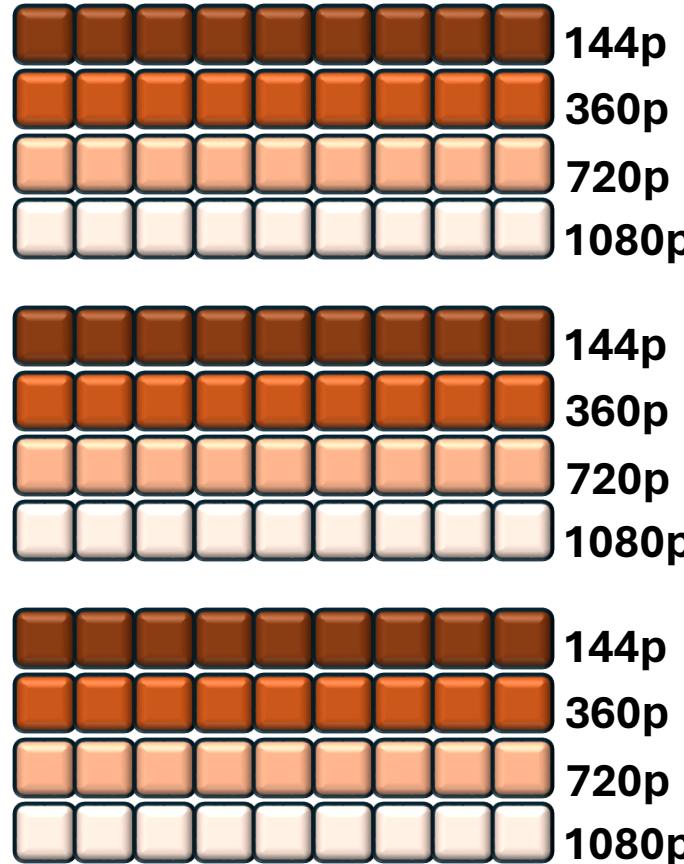
Is the next video ready for playing
when the user swipes?

Video Streaming Applications



Download next video

Download the next chunk of current video



Content distribution networks (CDNs)

challenge: how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?

- *option 1:* single, large “mega-server”
 - single point of failure
 - point of network congestion
 - long (and possibly congested) path to distant clients

....quite simply: this solution *doesn't scale*

Content distribution networks (CDNs)

challenge: how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?

- *option 2:* store/serve multiple copies of videos at multiple geographically distributed sites (*CDN*)
 - *enter deep:* push CDN servers deep into many access networks
 - close to users
 - Akamai: 240,000 servers deployed in > 120 countries (2015)
 - *bring home:* smaller number (10's) of larger clusters in POPs near access nets
 - used by Limelight



Akamai today:



The Akamai Edge Today

360K
servers

100+
million hits
per second

7+
trillion
deliveries
per day

175+
terabits per
second
(250+ peak)

4,200+
locations

1,350+
networks

840+
cities

135
countries

Source: <https://networkingchannel.eu/living-on-the-edge-for-a-quarter-century-an-akamai-retrospective-downloads/>

How does Netflix work?

- Netflix: stores copies of content (e.g., MADMEN) at its (worldwide) OpenConnect CDN nodes
- subscriber requests content, service provider returns manifest
 - using manifest, client retrieves content at highest supportable rate
 - may choose different rate or copy if network path congested

